

SDM4 in R: Analysis of Variance (Chapter 26)

Nicholas Horton (nhorton@amherst.edu)

April 19, 2016

Introduction and background

This document is intended to help describe how to undertake analyses introduced as examples in the Fourth Edition of *Stats: Data and Models* (2014) by De Veaux, Velleman, and Bock. More information about the book can be found at http://wps.aw.com/aw_deveaux_stats_series. This file as well as the associated R Markdown reproducible analysis source file used to create it can be found at <http://www.amherst.edu/~nhorton/sdm4>.


This work leverages initiatives undertaken by Project MOSAIC (<http://www.mosaic-web.org>), an NSF-funded effort to improve the teaching of statistics, calculus, science and computing in the undergraduate curriculum. In particular, we utilize the `mosaic` package, which was written to simplify the use of R for introductory statistics courses. A short summary of the R needed to teach introductory statistics can be found in the `mosaic` package vignettes (<http://cran.r-project.org/web/packages/mosaic>).

Chapter 26: Analysis of variance

Section 26.1: Testing whether the means of several groups are zero


The graph in Figure 26.1 (page 747) can be generated using the `bwplot()` function.

```
Soap <- read.csv("http://www.amherst.edu/~nhorton/sdm4/data/Bacterial_Soap.csv")
bwplot(Bacterial.Counts ~ Method, data=Soap)
```



The example on page 750 considers the outcomes in hand volumes for three treatments post surgery.

```
Contrast <- read.csv("http://www.amherst.edu/~nhorton/sdm4/data/Contrast_baths.csv")
bwplot(Hand.Vol.Chg ~ Treatment, data=Contrast)
```



The summary statistics at the bottom of page 751 can be calculated using `favstats()`.

```
favstats(Bacterial.Counts ~ Method, data=Soap)
```

```
##           Method min    Q1 median    Q3 max  mean    sd n missing
## 1   Alcohol Spray  5 17.75  34.5  52.75  82  37.5 26.560 8      0
## 2 Antibacterial Soap 20 72.25  91.5 113.00 164  92.5 41.963 8      0
## 3           Soap  51 79.75 105.0 112.25 207 106.0 46.959 8      0
## 4           Water  74 98.25 114.5 136.00 170 117.0 31.131 8      0
```

Section 26.2: The ANOVA table


The `aov()` function can be used to fit an analysis of variance model.

```
aovmod <- aov(Bacterial.Counts ~ Method, data=Soap)
summary(aovmod)
```

```
##           Df Sum Sq Mean Sq F value Pr(>F)
## Method      3  29882    9961    7.06 0.0011 **
## Residuals   28  39484    1410
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

This model has 3 degrees of freedom for the model (numerator) and 28 degrees of freedom for the error (denominator). The `xpf()` function can replicate the calculation of the exact p-value (and generate Figure 26.4, page 754).

```
xpf(7.0636, df1=3, df2=28)
```



```
## [1] 0.99889
```

The treatment means can be generated using `model.tables()` (see page 757).

```
model.tables(aovmod)
```

```
## Tables of effects
##
## Method
## Method
##      Alcohol Spray Antibacterial Soap      Soap
##      -50.75      4.25      17.75
##      Water
##      28.75
```

The residual standard deviation can be calculated (page 759).

```
n <- 32; k <- 4
sp <- sqrt(sum(resid(aovmod)^2/(n-k))); sp
```

```
## [1] 37.552
```

```
sqrt(1410)
```

```
## [1] 37.55
```

We can also see how the results are equivalent when fitting a regression model with indicators.

```
lmmod <- lm(Bacterial.Counts ~ Method, data=Soap)
msummary(lmmod)
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      37.5      13.3    2.82 0.00863 **
## MethodAntibacterial Soap    55.0      18.8    2.93 0.00669 **
## MethodSoap          68.5      18.8    3.65 0.00107 **
## MethodWater         79.5      18.8    4.23 0.00022 ***
##
## Residual standard error: 37.6 on 28 degrees of freedom
## Multiple R-squared:  0.431, Adjusted R-squared:  0.37
## F-statistic: 7.06 on 3 and 28 DF, p-value: 0.00111
```